PRODUCT BRIEF

Pilot is the *wisdom engine*.

Pilot replaces the content management system. It ingests what you know, shapes it through editorial intelligence, and projects it as wisdom — into any surface, for any audience, at any moment.

David Habib · February 2026 · Confidential

THE PROBLEM

The CMS was built for a world that no longer exists

Three pillars of the online content economy are collapsing simultaneously. Search traffic is evaporating — 60% of Google searches now end without a click. Advertising economics are fragmenting as privacy regulation dismantles behavioral targeting. Content production costs are collapsing as AI commoditizes artifact-level output.

The content management system has no answer. It was designed to manage authored artifacts — articles, pages, posts — in a pull-based world where consumers sought content out. That world is being replaced by a push-based paradigm where algorithms decide what gets seen, content must be format-fluid, and the value lies not in any individual artifact but in accumulated institutional wisdom.

> *Your CMS manages articles. Pilot cultivates wisdom.*

THE PRODUCT

What Pilot does

Pilot is a SaaS platform with a fundamentally different architecture. There are no content types, no schemas, no taxonomies, no templates. Knowledge enters in its natural form — a document, a photograph, a voice memo, a data feed — and is decomposed into a continuous semantic manifold. An editorial constitution shapes how that knowledge is projected to the world. Every output is nondeterministic, format-fluid, and traceable back to its human sources.

INGESTION

Any format, any source

Upload a Word doc, photograph a whiteboard, record a voice memo, connect an RSS feed. Multimodal AI normalizes everything into semantic embeddings. No schemas. No content types. No formatting requirements.

MANIFOLD

Wisdom, not articles

Knowledge lives in a continuous semantic space — a vector database paired with a knowledge graph. Deduplication, contradiction detection, and freshness scoring happen automatically. Every assertion traces back to its human source.

CONSTITUTION

Editorial intelligence

Ten tunable dials define your entity's editorial personality — tone, framing, selection bias, engagement-quality tension, boundary conditions. The editor

PROJECTION

Any surface, any moment

Web pages, conversational AI, social posts, newsletters, data feeds, documents. Each projection is shaped by the constitution, informed by the

authors the entity's disposition. The AI executes it at scale.

manifold, and tailored to the audience. Content is format-fluid by nature.

PROVENANCE

Every claim, traceable

Every projection carries metadata linking back through the knowledge graph to its contributing humans. Provenance is not an add-on — it is woven into every layer. Ask "where did this come from?" and get a definitive answer.

BYOK

Your AI, your costs

Bring your own API keys for OpenAI, Anthropic, Google, or any compatible provider. Pilot's value is the platform intelligence, not reselling tokens. You control your AI spend. As model costs fall, you benefit directly.

CONTENT MODEL

Five categories, zero schemas

Every piece of information in Pilot belongs to one of five behavioral categories. These are not content types — there are no fields, no templates, no taxonomies. They determine how the system stores, governs, and projects each piece of content.

CATEGORY	WHAT IT IS	HOW IT'S GENERATED
Artifact	Human-authored knowledge. Source material for the manifold. Stored immutably.	Human upload
Boilerplate	Stable institutional content — privacy policies, about pages, legal. Editor-locked.	Human authored
Hook	SEO-optimized entry content. Stable URLs. The bridge between search and wisdom.	Deterministic projection

CATEGORY	WHAT IT IS	HOW IT'S GENERATED
Exploratory	Wisdom in any format. Nondeterministic. Personalized. Constitution-shaped.	Nondeterministic projection
Meta	Information about information — analytics, trends, telemetry. Drives editorial decisions.	System-generated

ARCHITECTURE

The bimodal frontend

Pilot's reference web application serves both paths simultaneously. Static content renders instantly from cache. Agentic content streams from the wisdom manifold via MCP (Model Context Protocol). The consumer sees a single, seamless experience that deepens as they engage.

Deterministic Path

`REST / GraphQL → CDN Cache`

Boilerplate — static generation, CDN-cached

Hook — incremental static regeneration, stable URLs

Public artifacts — immutable, content-addressed

TTFB < 100ms · LCP < 500ms

⇄

Agentic Path

`MCP over SSE → Streaming UI`

Exploratory — streamed, constitution-shaped, personalized

Research desk — conversational AI with the entity's voice

Inline provenance — citations embedded in streaming output

The hook-to-exploratory transition

A consumer arrives via search on a hook page — fast, cached, SEO-optimized. As they engage (scroll, dwell, click), the page progressively deepens into exploratory content

streamed from the manifold. The transition feels like the page is opening up, not redirecting. Search engines see the hook content. Humans discover the wisdom beneath it.

Graceful degradation

If the LLM is slow, the exploratory zone falls back to static related topics. If the LLM is down, hook content serves normally and the exploratory zone disappears. If the MCP server is down, the webapp runs in deterministic-only mode. The system never goes dark because a third-party API is unavailable. Boilerplate and hook content have zero LLM dependency.

FOR EDITORS

A command center, not a content calendar

The editor doesn't write copy or approve individual articles. The editor shapes the system's judgment — setting the entity's tone, framing, selection bias, and boundary conditions through ten constitutional dials.

At the editor's fingertips

New artifacts with automated quality assessment. Sentiment and trend data from internal analytics and external feeds. AI-interpreted insights about reach, consumption, and audience behavior — not raw dashboards, but actionable summaries.

What the editor can do

Accept AI-suggested projections — the system surfaces opportunities: trending topics, manifold gaps, audience demand signals. **Chat with the manifold** — a conversational AI that has access to everything, configured as an active listener. **Turn the dials** — adjust mood, tone, urgency, neutrality, and watch the projections shift in real time.

GUARDRAILS

Built-in safety and accountability

🔒 **Quarantine by default**

Suspect artifacts — factually inconsistent, from new contributors, flagged by content analysis, or bulk-imported — are isolated from the manifold until an editor releases them. Quarantine means the embeddings don't exist, not just that access is restricted.

🔗 **Provenance is mandatory**

Every assertion traces back through the knowledge graph to its human source. Every projection carries attribution metadata. Ask "who said this?" and get an answer. Always.

📜 **Constitutional governance**

The editorial constitution is not a suggestion — it is an enforcement mechanism. Every projection is evaluated against the constitution's principles and boundary conditions. Hard constraints are never overridden by optimization signals.

🔑 **BYOK key hygiene**

API keys live in a secrets vault. Retrieved per-call, held in memory only during execution, never logged, never in error messages, never visible to platform staff. Key rotation is zero-downtime.

👁️ **Full audit trail**

Every constitutional change, every quarantine decision, every projection event is logged in the knowledge graph and exportable. The auditor can reconstruct the system's state at any point in time.

🛡️ **No raw HTML from AI**

The LLM never outputs HTML. It outputs structured component references from a pre-approved registry. The security boundary is architectural, not a sanitization filter.

🌐 **Personalization isolation**

Consumer preference vectors are per-consumer, encrypted at rest, and never included in shared cache keys. One consumer's personalization cannot leak into another's experience.

Three filters for every decision

1

Cost

A solo creator with a $20/month LLM budget can run a useful Pilot instance. Never call the LLM when you don't have to. Cache aggressively. Batch embeddings. Defer non-critical processing. Every feature has a cost estimate.

2

Modularity

Everything beyond core manifold operations is a plugin. LLM providers, ingestion sources, projection surfaces, analytics — all plugins. Core plugins follow the same interface as third-party plugins. No backdoors.

3

Simplicity

Prefer boring technology. One way to do things. No premature abstraction. The novelty is the product concept, not the infrastructure. Fewer than ten services. Under $50/month for a small tenant.

BUSINESS MODEL

Low barrier, high ceiling

Pilot's base subscription covers the platform — manifold storage, editorial constitution, projection logic, plugin framework, core plugins. Customers scale costs in two dimensions: manifold volume (storage) and inference spend (paid directly to their LLM provider, not to Pilot).

Revenue comes from base subscriptions (tiered by manifold volume), premium plugins, marketplace revenue share from third-party plugin sales, and enterprise features for media companies and regulated industries.

Pilot does not compete with WordPress, Contentful, or Sanity. It competes with the category itself. The competitive moat is architectural — no CMS can evolve into a wisdom engine without abandoning its schema-based foundation, which means abandoning its entire ecosystem.

FOUNDATIONS

Built on original research

Pilot is grounded in the frameworks developed across two books by David Habib:

Latent Vector traces how information storage technologies — from clay tablets to filing cabinets to databases — create the organizational structures that surround them. The CMS is the publishing industry's version of this pattern: the schema created content teams, editorial workflows, and an Information-Commercial Complex of template vendors and SEO consultants. Pilot eliminates the schema entirely.

About 5 Years documents the pull-to-push revolution in content distribution — the simultaneous collapse of search traffic, advertising economics, and content production costs. Pilot is the architecture built for the push-based world: format-fluid, algorithm-ready, and optimized for premium perspective rather than commodity information.

David Habib
Author of *Latent Vector* and *About 5 Years*

Confidential · February 2026